

[Notes Click](#)

Practice Problems of Machine Learning

Note: ETE is based on your entire syllabus. Here are some practice questions on after MTE syllabus

SL No	Questions
1	<p>Discuss about Roulette Wheel Selection algorithm and Rank based Selection?</p> <p><u>Roulette Wheel Selection</u>¹²³⁴:</p> <ul style="list-style-type: none">• This method is also known as fitness proportionate selection.• It is a genetic operator used for selecting potentially useful solutions for recombination.• In this method, every individual can become a parent with a probability which is proportional to its fitness. Therefore, fitter individuals have a higher chance of mating and propagating their features to the next generation.• The wheel is divided into n pies, where n is the number of individuals in the population. Each individual gets a portion of the circle which is proportional to its fitness value.• A fixed point is chosen on the wheel circumference and the wheel is rotated. The region of the wheel which comes in front of the fixed point is chosen as the parent. For the second parent, the same process is repeated.• It is clear that a fitter individual has a greater pie on the wheel and therefore a greater chance of landing in front of the fixed point when the wheel is rotated. Therefore, the probability of choosing an individual depends directly on its fitness. <p><u>Rank Based Selection</u>⁵²⁶⁷:</p> <ul style="list-style-type: none">• Rank selection first ranks the population and then every chromosome receives fitness from this ranking.• The worst will have fitness 1, second worst 2 etc. and the best will have fitness N (number of chromosomes in population). After this all the chromosomes have a chance to be selected.• This method can avoid premature convergence, which is an undesirable condition in a genetic algorithm where one extremely fit solution takes over the entire population in a few generations, leading to a loss of diversity.• It can be computationally expensive because it sorts the populations based on fitness value.
2	<p>Explain how locally weighted regression is different from linear regression</p> <p><u>Linear Regression</u>¹:</p> <ul style="list-style-type: none">• Linear Regression is a supervised learning algorithm used for computing linear relationships between input (X) and output (Y).• It learns a fixed set of parameters during the training phase.• <u>Once the model is trained, the training set can be erased from the machine as the model has already learned the required parameters</u>¹.• <u>However, this algorithm cannot be used for making predictions when there exists a non-linear relationship between X and Y</u>¹. <p><u>Locally Weighted Regression (LWR)</u>¹²:</p> <ul style="list-style-type: none">• <u>LWR is a non-parametric algorithm, meaning the model does not learn a fixed set of parameters as is done in ordinary linear regression</u>¹.• <u>Instead, parameters are computed individually for each query point</u>¹.

	<ul style="list-style-type: none"> • <u>While computing, a higher “preference” is given to the points in the training set lying in the vicinity of the query point than the points lying far away from it¹.</u> • <u>The weights fall exponentially as the distance between the query point and the training points increases¹.</u> • <u>There exists no training phase. All the work is done during the testing phase/while making predictions¹.</u> • <u>The dataset must always be available for predictions¹.</u> • <u>LWR methods are a generalization of k-Nearest Neighbour¹.</u> • <u>In terms of Mean Squared Error (MSE), LWR outperforms traditional Linear Regression³.</u>
3	<p>Discuss about Markov Decision Process (MDP) and Q-Learning.</p> <p>Markov Decision Process (MDP)¹²³⁴:</p> <ul style="list-style-type: none"> • <u>A Markov Decision Process (MDP) is a mathematical framework used for modeling decision making in situations where outcomes are partly random and partly under the control of a decision maker².</u> • <u>An MDP model contains:</u> <ul style="list-style-type: none"> ◦ <u>A set of possible world states S^1.</u> ◦ <u>A set of possible actions A^1.</u> ◦ <u>A real-valued reward function $R(s,a)^1$.</u> ◦ <u>A policy, which is the solution of the Markov Decision Process¹.</u> • <u>The Markov property states that the effects of an action taken in a state depend only on that state and not on the prior history¹.</u> • <u>MDPs are useful for studying optimization problems solved via dynamic programming².</u> <p>Q-Learning⁵⁶⁷⁸:</p> <ul style="list-style-type: none"> • <u>Q-learning is a model-free reinforcement learning algorithm used to learn the value of an action in a particular state⁵.</u> • <u>It does not require a model of the environment (hence “model-free”), and it can handle problems with stochastic transitions and rewards without requiring adaptations⁵.</u> • <u>For any finite Markov decision process, Q-learning finds an optimal policy in the sense of maximizing the expected value of the total reward over any and all successive steps, starting from the current state⁵.</u> • <u>“Q” refers to the function that the algorithm computes – the expected rewards for an action taken in a given state⁵.</u> • <u>Q-learning can identify an optimal action-selection policy for any given finite Markov decision process, given infinite exploration time and a partly random policy⁵.</u>
4	<p>Explain Bagging and boosting</p> <p>Bagging: Bagging, which stands for Bootstrap Aggregating, is a way to decrease the variance of your prediction by generating additional data for training from your original dataset using combinations with repetitions to produce multi-sets of the original data. It works as follows:</p> <ol style="list-style-type: none"> 1. Multiple subsets are created from the original dataset, selecting observations with replacement. 2. A base model (weak model) is created on each of these subsets. 3. The models run in parallel and are independent of each other. 4. The final predictions are determined by combining the predictions from all the models. <p>Boosting: Boosting is an iterative technique that adjusts the weight of an observation based on the last classification. If an observation was classified incorrectly, it tries to increase the weight of this observation and vice versa. Boosting in general decreases the bias error and builds strong</p>

	<p>predictive models. Here's how it works:</p> <ol style="list-style-type: none"> 1. A subset is created from the original dataset. 2. Initially, all data points are given equal weights. 3. A base model is created on this subset. 4. This model is used to make predictions on the whole dataset. 5. Errors are calculated using the actual values and predicted values. 6. The observations which are incorrectly predicted are given higher weights. 7. Another model is created and predictions are made on the dataset. (This model tries to correct the errors from the previous model) 8. Similarly, multiple models are created, each correcting the errors of the previous model. 9. The final model (strong learner) is the weighted mean of all the models (weak learners).
5	<p>Evaluate the performance of a Case-Based Learning system with an suitable example. What are the parameters to assess the effectiveness of case based learning system</p> <p>Case-Based Learning (CBL) is a method where students apply their knowledge to real-world scenarios. To evaluate a CBL system, you can look at:</p> <p><u>Case-Based Learning (CBL) is a pedagogical method where students apply their knowledge to real-world scenarios, promoting higher levels of cognition¹. It's often used in health professional education, where learning activities are commonly based on patient cases².</u></p> <p>To evaluate the performance of a CBL system, let's consider an example from medical education. In this scenario, students are presented with a patient case involving specific symptoms and medical history. The students must then apply their theoretical knowledge and reasoning abilities to diagnose the patient and propose a treatment plan. <u>The effectiveness of the CBL system can be evaluated based on the accuracy of the students' diagnoses and treatment plans, as well as their ability to justify their decisions using their theoretical knowledge</u></p> <ol style="list-style-type: none"> 1. Academic Performance: Measured by exam scores or other assessments. 2. Interest and Motivation: CBL can enhance students' motivation to learn. 3. Ability to Analyze Cases: CBL should increase students' ability to apply theoretical knowledge to real-world scenarios. 4. Integration of Learning: CBL facilitates the integration of learning, allowing students to connect theoretical knowledge with practical applications. 5. Self-Reflection and Critical Reflection: CBL encourages learners to reflect on their learning process and critically evaluate their own understanding. 6. Collaborative Learning: CBL utilizes collaborative learning, promoting teamwork and communication skills.
6	<p>Discuss different components of Genetic algorithm</p> <p>A Genetic Algorithm (GA) is an adaptive heuristic search algorithm inspired by the process of natural selection. Here are the key components of a GA:</p> <ol style="list-style-type: none"> 1. Population: This is a subset of all possible solutions to the given problem. <u>Each individual in the population represents a potential solution¹.</u> 2. Chromosomes: These are the solutions in the population for the given problem. <u>A chromosome is a string of genes, where each gene represents a variable component of the solution²¹.</u> 3. Genes: These are the elements of a chromosome. <u>A chromosome is divided into different genes¹.</u>

	<ol style="list-style-type: none"> 4. Allele: This is the value provided to a gene within a particular chromosome¹. 5. Fitness Function: This function is used to determine the fitness level of individuals in the population. It measures the ability of an individual to compete with other individuals. <u>Individuals with better fitness scores are given more chances to reproduce²¹.</u> 6. Genetic Operators: These operators play a role in changing the genetic composition of the next generation. <u>They include selection, crossover (recombination), and mutation²¹.</u> 7. Selection: <u>After calculating the fitness of every individual in the population, a selection process is used to determine which individuals will get to reproduce and produce the offspring that will form the next generation¹.</u> 8. Crossover: <u>This is the process of producing the next generation of chromosomes by combining the genes of parent chromosomes².</u> 9. Mutation: <u>This introduces small random changes in the offspring, which helps maintain genetic diversity in the population and prevents premature convergence².</u> 10. Termination: <u>This is the condition that determines when to stop the algorithm. Common termination conditions include reaching a maximum number of generations, achieving a satisfactory fitness level, or not seeing any improvement for a certain number of generations¹.</u>
7	<p>Give an detail comparison among k-means, hirerchichal and density based clustering methods</p> <p>K-means Clustering:</p> <ul style="list-style-type: none"> It's a partitioning method that divides a dataset into a certain number (k) of clusters. It starts with random centroids and assigns each data point to the nearest centroid. It's suitable for large datasets and when there's a preliminary understanding of how many clusters the dataset should be segmented into. It's particularly effective when the dataset has a globular cluster shape. <p>Hierarchical Clustering:</p> <ul style="list-style-type: none"> It's a method that builds a hierarchy of clusters. It starts with each data point as its own cluster and merges clusters until all data points belong to a single cluster. It's suitable for small datasets. It's like a family tree, where each group of relatives is nested within a larger group. <p>Density-Based Clustering (DBSCAN):</p> <ul style="list-style-type: none"> It operates on a density-based approach. It groups together points that are packed closely together (points with many nearby neighbors), marking low-density regions that are far from any point as outliers. It's good when dealing with spatial data, noise, and when the clusters are in arbitrary shapes.
8	<p>Explain the various steps involved in Partitional clustering algorithm. Use this algorithm to develop a model for a real life problem.</p> <p>in a Partitional clustering algorithm, specifically the K-Means algorithm:</p> <ol style="list-style-type: none"> Initialization: Randomly assign K objects from the dataset as cluster centers. Assignment: Assign each object to the cluster that has the most similar mean value. Update: Recalculate the mean of each cluster with the updated values. Repeat: Repeat the assignment and update steps until no change occurs.

Now, let's apply this algorithm to a real-life problem: customer segmentation for a retail business.

1. **Initialization:** The retail business decides to segment its customers into 5 groups ($K=5$). It randomly selects 5 customers as the initial cluster centers.
2. **Assignment:** Each customer is assigned to the cluster whose center (the initially selected customer) they are most similar to. Similarity could be based on factors like purchasing habits, frequency of shopping, average spending, etc.
3. **Update:** The mean value of the purchasing habits, frequency of shopping, average spending, etc., for the customers in each cluster is calculated. This mean value becomes the new center of the cluster.
4. **Repeat:** Steps 2 and 3 are repeated until the customers' cluster assignments no longer change.

At the end of this process, the retail business has a segmentation of its customers into 5 groups. It can then tailor its marketing strategies to the unique characteristics of customers in each group, potentially improving its marketing effectiveness. For example, customers in a cluster characterized by high-frequency, low-spending shopping habits might respond best to frequent, small-scale promotions, while customers in a cluster characterized by low-frequency, high-spending shopping habits might respond best to infrequent, large-scale promotions.

9 State various types of artificial neural network with their advantages and disadvantages

Type of ANN	Advantages	Disadvantages
Feedforward Neural Networks	<u>Simplest type of ANN with information that moves in only one direction, forward, from the input nodes, through the hidden nodes (if any) and to the output nodes¹.</u>	<u>They are unable to learn temporal sequences due to their feedforward structure¹.</u>
Radial Basis Function Networks	<u>They can approximate any function and are good at dealing with noise¹.</u>	<u>They require a good knowledge of the spread parameter¹.</u>
Recurrent Neural Networks (RNNs)	<u>They are dynamic and have 'memory' which makes them powerful for tasks where context from earlier inputs is required¹.</u>	<u>Training RNNs can be quite difficult due to the problem of vanishing gradients¹.</u>
Long Short-Term Memory Networks (LSTMs)	<u>They are a special kind of RNN capable of learning long-term dependencies, which makes them effective for sequence prediction problems¹.</u>	<u>They can be quite complex and computationally intensive¹.</u>
Convolutional Neural Networks (CNNs)	<u>They are very effective for image and video processing tasks due to their ability to process spatial information¹.</u>	<u>They require a large amount of training data and computational resources¹.</u>

10

Illustrate the steps of Agglomerative Hierarchical Clustering and solve the following dataset.

Points:

A (12, 13)

B (15, 14)

C (19, 16)

D (18, 12)

E (17, 15)

[YouTube Reference](#)

Note AI answer

Step	Clusters
Initialization	{A}, {B}, {C}, {D}, {E}
1st Merge (A, B closest)	{AB}, {C}, {D}, {E}
2nd Merge (Example: AB, C closest)	{ABC}, {D}, {E}
3rd Merge (Example: ABC, D closest)	{ABCD}, {E}
4th Merge (All points)	{ABCDE}

This table shows the agglomerative hierarchical clustering process for your dataset. The actual merges may vary based on the distances between clusters. The final result is a dendrogram that shows the order in which points and clusters were merged. This dendrogram can be cut at different heights to form different numbers of clusters. For example, cutting it at a height that corresponds to 3 clusters might yield the clusters {AB}, {C}, {DE}.

11

Construct a detailed explanation of the simple model of an Artificial Neuron and its functions. Develop an understanding of the role of weights and biases in an artificial neuron. Analyze how weights and biases influence the neuron's response to input signals, and discuss the mechanisms by which they contribute to the overall behavior of the neuron

An artificial neuron is a mathematical function modeled after biological neurons. The neuron receives one or more inputs, applies weights to these inputs, and sums them to produce an output.

Weights and biases are two essential components of an artificial neuron:

- Weights are numerical values associated with the connections between neurons. They determine the strength of these connections and the influence that one neuron's output has on another neuron's input. During the training phase of a neural network, these weights are adjusted iteratively to minimize the difference between the network's predictions and the actual outcomes.
- Biases are constants associated with each neuron. Unlike weights, biases are not connected to specific inputs but are added to the neuron's output.

Weights and biases play a crucial role in determining the neuron's response to input signals:

- Weights increase the steepness of the activation function, which means they decide how fast the activation function will trigger. The weights associated with each input determine how much importance the network places on that input when making a decision.
- Biases are used to delay the triggering of the activation function. The bias can increase or

decrease a neuron's output.

The weights and biases develop how a neural network propels data flow forward through the network; this is called forward propagation. As inputs are transferred across neurons, weights and biases are applied to the inputs. The change in weights and biases during the training process fine-tunes the network's ability to make accurate predictions. This iterative learning process of neural networks, involving both forward and backward propagation, contributes to the overall behavior of the neuron.

12

Identify and classify the various methods used for dimensionality reduction. Choose one method and provide a detailed explanation of its working principles, advantages, and limitations.

Method	Description	Advantages	Disadvantages
Feature Selection Methods	Select a subset of the original features that are most relevant to the problem.	Can reduce dimensionality while retaining important features.	May lose some information as features are completely removed.
Matrix Factorization Methods (e.g., PCA)	Use linear algebra techniques to reduce dimensionality.	Can capture most of the variance in the data using fewer dimensions.	Assumes linear relationships among features.
Manifold Learning Methods (e.g., t-SNE)	Find a low-dimensional basis for describing high-dimensional data.	Can handle non-linear relationships among features.	Computationally intensive and may not preserve distances between clusters.
Autoencoder Methods	Use neural networks to learn efficient codings of input data.	Can capture complex patterns in the data.	Require a lot of data and computational resources.

Now, let's focus on **PCA**:

Principal Component Analysis (PCA) is a technique used to emphasize variation and bring out strong patterns in a dataset. It's often used to make data easy to explore and visualize.

Working Principle:

- PCA starts by identifying the hyperplane that lies closest to the data, and then it projects the data onto it.
- The direction (or principal component) with the largest variance is considered as the first principal component. The direction orthogonal to the first component and has the second largest variance is considered as the second principal component, and so on.

Advantages:

- PCA is not affected by scale, and you can use it on features with different scales.
- It can drastically reduce the dimension of the data while keeping most of the valuable

	<p>parts.</p> <p>Limitations:</p> <ul style="list-style-type: none"> • PCA assumes that the principal components are a linear combination of the original features. If this is not the case, PCA may not give you a sensible result. • It also treats the principal components as orthogonal. This is a limitation when dealing with correlated features.
13	<p>Develop a comprehensive analysis of the computational complexity of linear regression and nonlinear regression models. How does the complexity of the model affect the training and inference time? Discuss the key differences in terms of their assumptions, functional forms, and interpretability.</p> <p>Computational Complexity of Linear and Nonlinear Regression Models:</p> <ul style="list-style-type: none"> • Linear Regression: The computational complexity of linear regression is primarily determined by the number of features and the number of training examples. <u>Specifically, the complexity of matrix multiplication, which is a key operation in linear regression, is $O(C^2N)$ where N is the number of training examples and C is the total number of features¹.</u> • Nonlinear Regression: Nonlinear regression models can be more complex to implement than linear regression models. <u>They can be more sensitive to outliers and more computationally expensive to train².</u> <p>Effect of Model Complexity on Training and Inference Time:</p> <ul style="list-style-type: none"> • Training Time: <u>The resources required for training, such as memory, storage, and compute, increase with both the size of the dataset and the complexity of the model³. For example, in linear regression, it means solving the optimization problem, which has a time complexity of $O(nd)$, where n is the number of training examples and d is the number of dimensions⁴.</u> • Inference Time: <u>Once a model is trained, you can evaluate new instances within fractions of a second, so the training time is well invested⁵.</u> <p>Key Differences in Terms of Assumptions, Functional Forms, and Interpretability:</p> <ul style="list-style-type: none"> • Assumptions: <u>Linear regression assumes a linear relationship between the independent and dependent variables, while nonlinear regression does not make this assumption⁶².</u> • Functional Forms: <u>Linear regression models can model curvature by including nonlinear variables such as polynomials and transforming exponential functions⁷. Nonlinear regression models can accommodate different mean functions, even though they are less flexible than linear regression models⁸.</u> • Interpretability: <u>Linear regression models are generally easier to interpret than nonlinear models because the relationship between the independent and dependent variables is simply a straight line. Nonlinear models, on the other hand, can be more difficult to interpret due to their complexity⁹.</u> <p>In conclusion, both linear and nonlinear regression models have their own advantages and disadvantages, and the choice between them depends on the specific characteristics of the data and the problem at hand.</p>

14	<p>Differentiate between discriminative learning algorithms and generative learning algorithms.</p> <table> <tr> <th>Discriminative Learning Algorithms</th><th>Generative Learning Algorithms</th></tr> <tr> <td>Description</td><td> <p>These algorithms model the decision boundary between the classes. They are designed to learn mappings directly from the input space to the output space. They do not attempt to model how the data was generated; they simply categorize a given signal.</p> <p>These algorithms model how the data is distributed in each class. They learn the joint probability distribution of the input and output. After learning the distribution, generative models can generate new data points that reflect the training set.</p> </td></tr> <tr> <td>Examples</td><td> <p>Logistic regression, support vector machines, and most neural networks.</p> <p>Naive Bayes, Gaussian Mixture Models, and Hidden Markov Models.</p> </td></tr> <tr> <td>Advantages</td><td> <p>They generally perform better for classification tasks because they focus on solving the actual problem directly.</p> <p>They are useful for tasks like text generation, where the model picks up the grammar and creates logical text sequences.</p> </td></tr> <tr> <td>Disadvantages</td><td> <p>They may not perform well when the classes are not linearly separable or when the number of features is very large compared to the number of training examples.</p> <p>They can be more sensitive to outliers and more computationally expensive to train.</p> </td></tr> </table>	Discriminative Learning Algorithms	Generative Learning Algorithms	Description	<p>These algorithms model the decision boundary between the classes. They are designed to learn mappings directly from the input space to the output space. They do not attempt to model how the data was generated; they simply categorize a given signal.</p> <p>These algorithms model how the data is distributed in each class. They learn the joint probability distribution of the input and output. After learning the distribution, generative models can generate new data points that reflect the training set.</p>	Examples	<p>Logistic regression, support vector machines, and most neural networks.</p> <p>Naive Bayes, Gaussian Mixture Models, and Hidden Markov Models.</p>	Advantages	<p>They generally perform better for classification tasks because they focus on solving the actual problem directly.</p> <p>They are useful for tasks like text generation, where the model picks up the grammar and creates logical text sequences.</p>	Disadvantages	<p>They may not perform well when the classes are not linearly separable or when the number of features is very large compared to the number of training examples.</p> <p>They can be more sensitive to outliers and more computationally expensive to train.</p>
Discriminative Learning Algorithms	Generative Learning Algorithms										
Description	<p>These algorithms model the decision boundary between the classes. They are designed to learn mappings directly from the input space to the output space. They do not attempt to model how the data was generated; they simply categorize a given signal.</p> <p>These algorithms model how the data is distributed in each class. They learn the joint probability distribution of the input and output. After learning the distribution, generative models can generate new data points that reflect the training set.</p>										
Examples	<p>Logistic regression, support vector machines, and most neural networks.</p> <p>Naive Bayes, Gaussian Mixture Models, and Hidden Markov Models.</p>										
Advantages	<p>They generally perform better for classification tasks because they focus on solving the actual problem directly.</p> <p>They are useful for tasks like text generation, where the model picks up the grammar and creates logical text sequences.</p>										
Disadvantages	<p>They may not perform well when the classes are not linearly separable or when the number of features is very large compared to the number of training examples.</p> <p>They can be more sensitive to outliers and more computationally expensive to train.</p>										
15	<p>How Gaussian Discriminant Analysis is related to Logistic Regression? Describe it in detail.</p> <p>Gaussian Discriminant Analysis (GDA): GDA is a generative learning algorithm used for classification tasks. It works by assuming that the data in each class follows a Gaussian (normal) distribution, and then estimating the mean and covariance matrix for each class. It then uses Bayes' theorem to compute the probability that a new data point belongs to each class, and chooses the class with the highest probability as the predicted class.</p> <p>Logistic Regression: Logistic regression is a discriminative learning algorithm used for binary classification tasks. It predicts the probability of an outcome, such as yes/no, 0/1, or true/false. In logistic regression, a logit transformation is applied on the odds—that is, the probability of success divided by the probability of failure.</p> <p>Relation between GDA and Logistic Regression: Both GDA and Logistic Regression are</p>										

supervised learning algorithms used for classification. The key difference between them lies in the way they approach the problem:

- GDA is a generative model, which first models the distribution of individual classes and then uses Bayes' theorem to make predictions. It makes more specific assumptions about the data set than Logistic Regression and if those assumptions are true then it works better than LR.
- Logistic Regression, on the other hand, is a discriminative model, which directly learns the decision boundary between the classes. It makes more generic assumptions and can be more useful in a lot of other places where the probability distribution of the feature set is not Gaussian.

In summary, if the assumptions of GDA about the data are correct, it can outperform Logistic Regression. However, if these assumptions are violated, Logistic Regression can often be a more robust choice.

16 How does Singular Value Decomposition (SVD) work to decompose a matrix?

Singular Value Decomposition (SVD) is a factorization method that decomposes a matrix into three separate matrices. Given a matrix A , the SVD is given by the formula:

$$A = UDV^T$$

Here's how each component is defined:

- **U:** An $m \times m$ matrix where the columns are the orthonormal eigenvectors of AA^T .
- **D:** A diagonal matrix where the diagonal elements are the singular values of A . These singular values are the square roots of the eigenvalues of either AA^T or A^TA .
- **V^T :** The transpose of an $n \times n$ matrix where the columns are the orthonormal eigenvectors of A^TA .

The process of calculating the SVD of a matrix involves the following steps:

1. **Calculate the eigenvectors and eigenvalues of AA^T and A^TA :** These will be used to form the matrices U and V .
2. **Form the matrix U using the eigenvectors of AA^T :** Each column of U is an eigenvector of AA^T .
3. **Form the matrix V using the eigenvectors of A^TA :** Each column of V is an eigenvector of A^TA .
4. **Form the diagonal matrix D :** The elements along the diagonal are the square roots of the eigenvalues from either AA^T or A^TA . These are known as the singular values of A .

The resulting matrices U , D , and V^T give the singular value decomposition of the matrix A . This decomposition has many useful properties and forms the basis for many algorithms in numerical linear algebra, such as pseudoinverses and least squares solutions.

17 Examine the Expectation-Maximization (EM) algorithm and its iterative optimization method that combines various unsupervised machine learning algorithms

The Expectation-Maximization (EM) algorithm is an iterative method used in machine learning and statistics to find estimates of parameters in statistical models that involve latent variables. It's commonly used for latent variable models and can handle missing data.

The EM algorithm consists of two steps that are repeated iteratively: the Expectation step (E-step) and the Maximization step (M-step).

E-step: In this step, the algorithm computes the expectation of the log-likelihood using the current parameter estimates. This involves estimating the values for the latent variables.


M-step: In this step, the algorithm determines the parameters that maximize the expected log-likelihood obtained in the E-step. The model parameters are updated based on the estimated latent variables.

By iteratively repeating these steps, the EM algorithm seeks to maximize the likelihood of the observed data. It's commonly used for unsupervised learning tasks, such as clustering, where latent variables are inferred.

The EM algorithm combines various unsupervised machine learning algorithms by providing a general framework for learning in the presence of unobserved variables. It's particularly useful in situations where data is missing or incomplete. It serves as the foundation for many unsupervised clustering algorithms in the field of machine learning.

Suppose, we have 5 rooms in a building connected by doors as shown in the figure below. We shall number each room 0 through 4. The outside of the building can be thought of as one big room i.e. number 5. Notice that doors 1 and 4 lead into the building to room 5 (goal point). Show what will be the Q matrix after 2 episodes.


18



[YouTube Link Reference](#)

Consider the following Markov Random Field.

Justify your answer which of the following nodes will have no effect on D given the Markov Blanket of D.



19

[Utube](#)

The dataset of win or loss in a sport of 5 players is given in the table. Use logistic regression as classifier to answer the following questions:

Hours played	win(1)/ loss(0)
20	0
12	1
30	0
22	0
35	1

Calculate the probability of win for the player who played 39 hours.

20

[Link Reference](#)

Develop a solution using the Weighted k-NN algorithm to determine the class of a test instance (7.6, 60, 8) based on a provided training dataset in a table. Set the value of K to 3 and apply the algorithm to calculate the weighted class by considering the distances and weights of the nearest neighbors.

S.No.	CGPA	Assessment	Project Submitted	Result
1	9.2	85	8	Pass
2	8	80	7	Pass
3	8.5	81	8	Pass
4	6	45	5	Fail
5	6.5	50	4	Fail
6	8.2	72	7	Pass
7	5.8	38	5	Fail
8	8.9	91	9	Pass

[Link](#)

1. Calculate Euclidean distances between the test instance (7.6, 60, 8) and each instance in the training dataset:

- $D1 = \sqrt{(7.6-9.2)^2 + (60-85)^2 + (8-8)^2} = 25.9$
- $D2 = \sqrt{(7.6-9)^2 + (60-80)^2 + (8-7)^2} = 21$
- $D3 = \sqrt{(7.6-8.5)^2 + (60-81)^2 + (8-8)^2} = 21.4$
- $D4 = \sqrt{(7.6-6)^2 + (60-45)^2 + (8-5)^2} = 17$
- $D5 = \sqrt{(7.6-6.5)^2 + (60-50)^2 + (8-4)^2} = 15$
- $D6 = \sqrt{(7.6-8)^2 + (60-72)^2 + (8-7)^2} = 12$
- $D7 = \sqrt{(7.6-5.8)^2 + (60-33)^2 + (8-5)^2} = 27$

2. Select the K=3 nearest neighbors. The three closest neighbors are instances with distances of 12, 15, and 17.

3. Calculate weights using these distances:

- $W1 = 1/12$
- $W2 = 1/15$
- $W3 = 1/17$

4. Multiply these weights by their respective class labels:

- Pass: $W1 * \text{Pass}(1)$
- Fail: $W3 * \text{Fail}(0)$

5. Result: Pass(0.088), Fail(0.059)

Since "Pass" has a higher weighted sum than "Fail", we predict that the test instance belongs to class "Pass".

Note: The values in parentheses represent a binary encoding for "Pass" and "Fail", where Pass is encoded as '1' and Fail is encoded as '0'. The final class is determined by the class with the highest weighted sum. In this case, "Pass" has a higher weighted sum than "Fail", so the test instance is predicted to belong to the "Pass" class.

22 Discuss about deep Q-learning

Deep Q-Learning is a model-free reinforcement learning algorithm that combines the power of deep neural networks with the Q-learning algorithm¹²³. It enables an agent to learn how to make optimal decisions in an environment by maximizing cumulative rewards³.

Here's a step-by-step explanation of how Deep Q-Learning works:

1. **Initialize the Q-function:** The Q-function, which represents the expected cumulative reward of taking a certain action in a certain state and following a certain policy, is approximated using a deep neural network¹².

2. **Calculate Distances:** The agent interacts with the environment and collects experiences, which are stored in a replay buffer¹.
3. **Experience Replay:** Random batches of experiences are sampled from the replay buffer and used to train the network. This helps to break the correlation between consecutive experiences and stabilize the training process¹.
4. **Update the Q-function:** The Q-function is updated iteratively as the agent interacts with the environment¹.
5. **Policy Improvement:** The policy, which determines the agent's behavior, is improved based on the updated Q-values¹.

One of the key challenges in implementing Deep Q-Learning is that the Q-function is typically non-linear and can have many local minima. This can make it difficult for the neural network to converge to the correct Q-function¹. To address this, several techniques have been proposed, such as experience replay and target networks¹.

Deep Q-Learning has been used in various applications such as game playing, robotics, and autonomous vehicles¹. It's particularly useful in environments with a large number of states and actions, as well as for learning from high-dimensional inputs such as images or sensor data

